

# Optical Character Recognition

Stelios Barberakis  
<[chefarov@gmail.com](mailto:chefarov@gmail.com)>

Machine learning  
Chania, TUC  
26/03/2014



# Introduction

- What is *Optical Character Recognition (OCR)*?
- Why do we need OCR technology?
  - Usefull applications:
    - Assisting blind and/or visually impaired people
    - Automatic Postcode recognition
    - And more...
  - First small step towards defeating Captcha



# Introduction (2)


Evaluated some *Machine Learning (ML)* methods/techniques:

- Multilayer Perceptron Networks (MLP)
  - Both a two (2) and three (3) hidden layer
- K-Nearest Neighbors (knn)
- Support Vector Machine (SVM)

Working on the input space

# MNIST Dataset

- Used for Comparing methods
- Train set: 60000 sample
- Test set: 10000 samples
- Test set: people not participating in training set
- Resolution: 28x28
- Train set: random ordering
- Test set: 5000 easy, then 5000 harder





# Related Work

- MNIST test error rate:
  - Linear Classifiers: 7.6 - 12%
  - KNN: 0.63 - 5%
  - Boosted Stumps: 0.87 - 7.7%
  - PCA with Quadratic Classifiers: 3.3%
  - SVM: 0.56 - 1.4%
  - ANN (e.g. MLP, CNN): 0.23 – 4.7%
    - **MCDNN: 0.23%**
  - Human error: 0.2%

# Best implementation so far

- CNNs: state of the art classifiers for image recognition. Natively translational and scaling invariant
- MCDNN by Ciresan, Meier, Schmidhuber
- Test Error: 0.23%, Human error: 0.2%
- 20/23 of misclassified: correct 2<sup>nd</sup> guess
- Traffic signs dataset: Achieved half of the human error (1.1% → 0.54%)




# Preprocessing


- Removed padding and scaled pictures to 14x14
- Performance of scaled dataset similar to the one on the original dataset (10% used)
- Scaling intensity values to  $[-1, 1]$ . Pictures with different intensity ranges can be classified

*note: before calling the trained classifier for a random picture, it must be also normalized to  $[-1, 1]$*

# Multilayer Perceptron Network (MLP)



- Fully connected feed-forward network
- Non-linear classifier
- Input:  $M \times N$ ,  $M$  samples, each of  $N$  attributes
- Batch (native) or sequential training (faster- preferred for huge datasets).
- 3 layer with sigmoidal activation function can approximate any function to arbitrary accuracy





# Backpropagation (BP) – Gradient Descent (GD)

- Learning occurs by adjusting weights in order to decrease error:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{k=1}^N (t_k - y_k)^2$$

- Derivation show that in order to follow the gradient downwards we update weights by:

$$w_{ik} \leftarrow w_{ik} + \eta(t_k - y_k)x_i,$$

- Other BP training functions also exist



# Training procedure

## Forwards phase:

- compute the activation of each neuron  $j$  in the hidden layer(s) using:

$$h_j = \sum_i x_i v_{ij} \quad (3.4)$$

$$a_j = g(h_j) = \frac{1}{1 + \exp(-\beta h_j)} \quad (3.5)$$

- work through the network until you get to the output layers, which have activations:

$$h_k = \sum_j a_j w_{jk} \quad (3.6)$$

$$y_k = g(h_k) = \frac{1}{1 + \exp(-\beta h_k)} \quad (3.7)$$

## Backwards phase:

- compute the error at the output using:

$$\delta_{ok} = (t_k - y_k) y_k (1 - y_k) \quad (3.8)$$

- compute the error in the hidden layer(s) using:

$$\delta_{hj} = a_j (1 - a_j) \sum_k w_{jk} \delta_{ok} \quad (3.9)$$

- update the output layer weights using:

$$w_{jk} \leftarrow w_{jk} + \eta \delta_{ok} a_j^{\text{hidden}} \quad (3.10)$$

- update the hidden layer weights using:

$$v_{ij} \leftarrow v_{ij} + \eta \delta_{hj} x_i \quad (3.11)$$

randomise the order of the input vectors so that you don't train in exactly the same order each iteration



# Stopping

To avoid overfitting: Earlystopping or L1,L2 regularization (extra term to the loss function)

- EarlyStopping:
  - Iteratively call train(for some epochs). Use validation set to determine how well the network has generalized.
  - Heuristics to decide when to stop.
  - Simplest one: Stop if validation error in last 2 or 3 (earlystopping-iterations) was increasing

# Parameters

- Number of layers (2 is usually enough)
- Number of nodes in each layer
- Activation function (each node): sigmoid function: differentiable. Neuron behavior. Gradient descent
- Learning rate( $\eta$ ): How fast it converges. The lower the safer (Reduce oscillations, and avoiding local minimum)
- Momentum( $\alpha$ ):  $\Delta w$  depends on the previous  $\Delta w$ s. Allows lower ( $\eta$ ) and increases convergence speed.

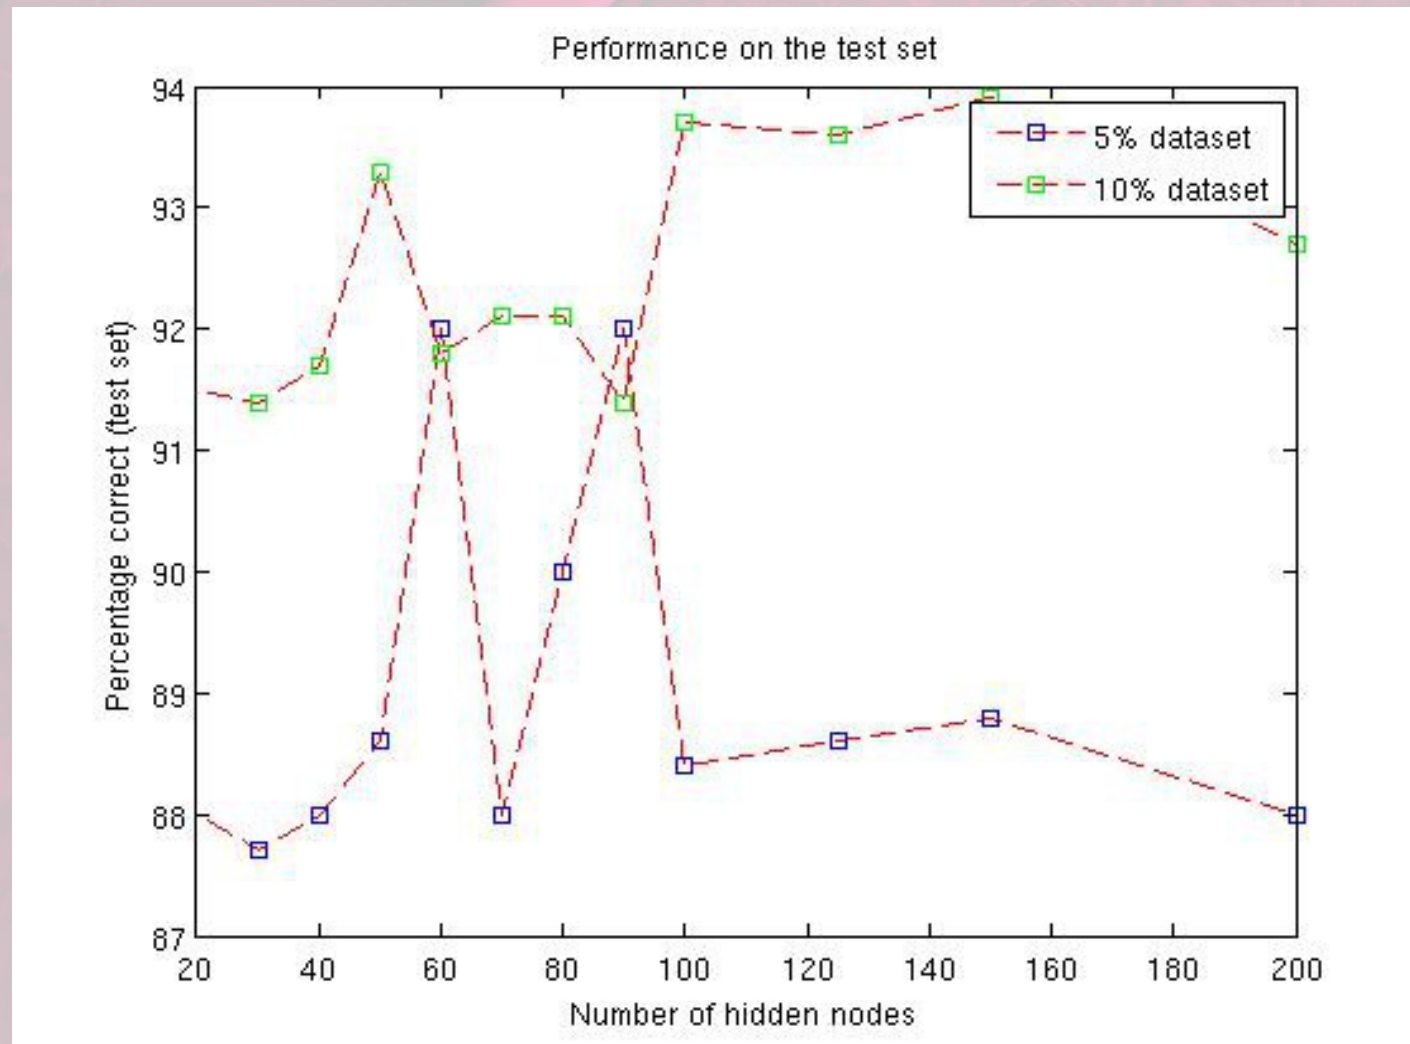
$$w_{ij}^t \leftarrow w_{ij}^{t-1} + \eta \delta_o a_j^{\text{hidden}} + \alpha \Delta w_{ij}^{t-1},$$



# Applying on MNIST dataset

- Input: Arrays (50000x14x14), reshape to (50000x196)
- Construct MLP (X, T, nodes, layers, outtype)
- Train(X,T,  $\eta$ , n\_iterations) #before calling earlystopping!
- Earlystopping(X,T,valid\_X,valid\_T,  $\eta$ )
- Calculate confusion matrix (test\_X, test\_T)
- Save trained NN (serialized in disk) using Pickle
- Recall phase: Call forward function for custom input
- Output: 1-of-N encoding using soft-max:  
[0.02, 0.30, 0.14 ... 0.05]: classified as '1' (argmax=1)

# Performance per architecture (2 layer - MLP)





# Performance of MLP

Architecture (nodes)	Dataset usage %	Performance	Training time
50	50	95.60	14h 40m
50	100	96.83	33h 20m
100	100	97.44 (vs 95.3 of 300-nodes reported)	66h 0m
-----	-----	-----	-----
150-100	100	97.17	5days 15h
200-100	100	97.41	52h
300-100	100	97.64 (best reported: 97)	8days

- Using Rprop from Pybrain (50 nodes) gave: 99.1% performance !?

# Optimize performance

- Parallelize Back propagation: Each thread calculates activation function for each neuron(sequential learning), or dataset is divided in subsets (batch learning).
- Use other training functions instead of Gradient Descent, as is the resilient Back Propagation (Rprop):
  - Each weight has its own update rule in the opposite direction of that weight's partial derivative.
- Best (paranoic) MLP reported:2500-2000-1500-1000-500-10, has performance: 98.43%, affine + elastic preprocessing: 99.69%



# Advantages & Disadvantages

- + Invariant features detection
- + General applicability
- + Can approximate complex models
- - A lot of experimentation needed
- - Training phase too long, need for parallelization
- - Difficult to make use of prior knowledge

# Script to classify a character

Parameters:

1. Image path
2. trained net object path
3. Foreground color



```
chefarov@debian:~/programming/ML/ocr14$ python classify.py -p samples/7e.jpg -n NN_s_100_0.save
[[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  1  0  0]
 [ 0  0  0  10  5  2  2  2  2  2  2  0  0  0]
 [ 0  1  0  37  36  31  38  42  48  57  54 101 16  0]
 [ 0  2  0  72 153 142 151 155 128  99 109 199 18  0]
 [ 0  0  0  14  28  24  6  44 164 160 197 124  0  4]
 [ 0  0  0  0  0  0  98 153 209 220 133 14  1  1]
 [ 0  0  0  1  0  82 214 199 219  70  0  0  0  0]
 [ 0  1  2  6 134 197 175 205  73  0  5  1  0  0]
 [ 0  2  0  95 202 179 191 155  0  4  0  0  0  0]
 [ 1  0 36 169 173 169 207  66  1  3  0  0  0  0]
 [ 1  0 39 167 171 180 186  7  0  1  0  0  0  0]
 [ 0  0  0 18  20 19 14  2  0  0  0  0  0  0]
 [ 0  0  1  0  0  0  0  0  0  0  0  0  0  0]]
[[ 1.04696293e-04  1.19935424e-01  3.33057029e-04  1.86990869e-03
  3.31822519e-06  2.56551445e-08  1.47507321e-07  6.99006362e-01
  1.78744038e-01  2.87642129e-06]]
The number was classified as 7
```





# kNN

- Non-parametric model, predictions straight from the data. No need for training phase
- For each test point to be classified:
  - ✓ Compute distance (usually euclidean), to every candidate
  - ✓ Predict by averaging the closets k elements.
- Use cross validation to choose k. Not necessary when large datasets are available
- We found out  $k=3$  as the best value for small subsets of the dataset

# Performance

- Using the entire dataset we got : 96.4% accuracy (3h 37m), working in input space
- Best performance reported (without preprocessing/FE): 96.91%
- Feature extraction boosts the performance by making the model robust against various input space distortions like (in OCR): translation, rotation, scaling, thinning, etc
- Best results achieved by
  - Shape Contexts (reported 99.37% performance)
  - Deformation models (reported: up to 99.48% performance)



# Conclusion on knn

- + Very simple, works surprisingly well in various scenarios
- + Easily implemented and parameterized
- + Useful for quick experiments
- - Doesn't learn from the data. Repeats the  $N$  comparisons and sorting for each classification
- - Sensitive to noisy data
- - Doesn't generalize well

# SVMs

- Binary classifier. Using multiple binary classifiers (eg: one-versus-one:  $n(n-1)/2$ ) to perform multiclass classification
- Map the input in higher dimensional space in order to be linearly separable  $\Phi(x)$
- Find a linear separating hyperplane with the maximal margin
- Given  $(x_i, y_i)$ ,  $i=1, \dots, l$

where  $x_i \in \mathbb{R}^n$

and  $y \in \{1, -1\}$

optimize:

$$\begin{aligned} \min_{\mathbf{w}, b, \xi} \quad & \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^l \xi_i \\ \text{subject to} \quad & y_i (\mathbf{w}^T \phi(\mathbf{x}_i) + b) \geq 1 - \xi_i, \\ & \xi_i \geq 0. \end{aligned}$$



# SVMs - steps

- Transform data to the format of an SVM package
- Conduct simple scaling on the data
- Consider the RBF kernel
- Use cross-validation on a small subset to find a good neighborhood for  $(C, \gamma)$ .
- Make a finer grid search on the neighborhood to find better approximation for the optimal  $(C, \gamma)$
- Use the best parameter  $C$  and  $\gamma$  to train the whole training set
- Evaluate generalization performance on the test set

$$K(\mathbf{x}, \mathbf{y}) = e^{-\gamma \|\mathbf{x} - \mathbf{y}\|^2}$$

# SVMs – basic kernels

- Polynomial more complex more parameters.
- Linear kernels only when data linearly separable

- linear:  $K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$ .
- polynomial:  $K(\mathbf{x}_i, \mathbf{x}_j) = (\gamma \mathbf{x}_i^T \mathbf{x}_j + r)^d, \gamma > 0$ .
- radial basis function (RBF):  $K(\mathbf{x}_i, \mathbf{x}_j) = \exp(-\gamma \|\mathbf{x}_i - \mathbf{x}_j\|^2), \gamma > 0$ .
- sigmoid:  $K(\mathbf{x}_i, \mathbf{x}_j) = \tanh(\gamma \mathbf{x}_i^T \mathbf{x}_j + r)$ .

- using 10% dataset, performed grid search (each parameter 6-7 values), for poly and rbf kernel's parameters:
- RBF kernel: 50min, Polynomial kernel: 2h



# performance

Making a grid search on MNIST gave as an optimal

- RBF:  $(C, \gamma)$ :  $C=2.8$ ,  $\gamma=0.0073$ . Performance: 98.6%,  
Same performance like the MNIST page reports
- Poly  $(C, \gamma, \rho, d)$ :  $C=0.35$ ,  $\rho=0.125$ ,  $\gamma=0.0625$ ,  $d=3$ .  
Performance: 98,25%, some minutes
- Stronger implementations:
  - Poly: Degree:4 reported 98.9% performance
  - Best: Virtual SVM, deg-9 poly, 2-pixel jittered  
achieved 99.44%

# Conclusion – Future work

- Deeply researched issue – human competitive results
- From individual character recognition to more general tasks
- I don't intent to further study this issue...
- If anyone comes up with this, consider using CNN's